

Data__Science__Capstone__HealthCare__Project__By__Keshav

October 1, 2022

```
[77]: #import the library
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
[2]: #Load the dataset
from google.colab import files
uploaded = files.upload()
```

<IPython.core.display.HTML object>

Saving health care diabetes.csv to health care diabetes.csv

```
[3]: data = pd.read_csv('health care diabetes.csv')
```

```
[4]: data.head()
```

```
[4]:
```

| | Pregnancies | Glucose | BloodPressure | SkinThickness | Insulin | BMI | \ |
|---|-------------|---------|---------------|---------------|---------|------|---|
| 0 | 6 | 148 | 72 | 35 | 0 | 33.6 | |
| 1 | 1 | 85 | 66 | 29 | 0 | 26.6 | |
| 2 | 8 | 183 | 64 | 0 | 0 | 23.3 | |
| 3 | 1 | 89 | 66 | 23 | 94 | 28.1 | |
| 4 | 0 | 137 | 40 | 35 | 168 | 43.1 | |

| | DiabetesPedigreeFunction | Age | Outcome |
|---|--------------------------|-----|---------|
| 0 | 0.627 | 50 | 1 |
| 1 | 0.351 | 31 | 0 |
| 2 | 0.672 | 32 | 1 |
| 3 | 0.167 | 21 | 0 |
| 4 | 2.288 | 33 | 1 |

```
[5]: data.shape
```

```
[5]: (768, 9)
```

```
[6]: data.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Pregnancies                          768 non-null    int64
1   Glucose                             768 non-null    int64
2   BloodPressure                       768 non-null    int64
3   SkinThickness                       768 non-null    int64
4   Insulin                             768 non-null    int64
5   BMI                                 768 non-null    float64
6   DiabetesPedigreeFunction            768 non-null    float64
7   Age                                 768 non-null    int64
8   Outcome                             768 non-null    int64
dtypes: float64(2), int64(7)
memory usage: 54.1 KB

```

```
[7]: data.describe()
```

```

[7]:      Pregnancies      Glucose  BloodPressure  SkinThickness      Insulin  \
count      768.000000    768.000000      768.000000      768.000000    768.000000
mean         3.845052    120.894531        69.105469        20.536458     79.799479
std          3.369578     31.972618        19.355807        15.952218    115.244002
min           0.000000      0.000000         0.000000         0.000000      0.000000
25%           1.000000     99.000000        62.000000         0.000000      0.000000
50%           3.000000    117.000000        72.000000        23.000000     30.500000
75%           6.000000    140.250000        80.000000        32.000000    127.250000
max          17.000000    199.000000       122.000000        99.000000    846.000000

      BMI  DiabetesPedigreeFunction      Age      Outcome
count    768.000000              768.000000    768.000000    768.000000
mean      31.992578                0.471876     33.240885     0.348958
std        7.884160                0.331329     11.760232     0.476951
min         0.000000                0.078000     21.000000     0.000000
25%        27.300000                0.243750     24.000000     0.000000
50%        32.000000                0.372500     29.000000     0.000000
75%        36.600000                0.626250     41.000000     1.000000
max        67.100000                2.420000     81.000000     1.000000

```

```
[8]: data.head()
```

```

[8]:      Pregnancies  Glucose  BloodPressure  SkinThickness  Insulin  BMI  \
0              6      148              72              35         0  33.6
1              1       85              66              29         0  26.6
2              8      183              64              0         0  23.3
3              1       89              66              23        94  28.1
4              0      137              40              35       168  43.1

```

| | DiabetesPedigreeFunction | Age | Outcome |
|---|--------------------------|-----|---------|
| 0 | 0.627 | 50 | 1 |
| 1 | 0.351 | 31 | 0 |
| 2 | 0.672 | 32 | 1 |
| 3 | 0.167 | 21 | 0 |
| 4 | 2.288 | 33 | 1 |

```
[9]: data.tail()
```

```
[9]:      Pregnancies  Glucose  BloodPressure  SkinThickness  Insulin   BMI  \
763           10     101           76           48       180  32.9
764           2     122           70           27         0  36.8
765           5     121           72           23       112  26.2
766           1     126           60           0         0  30.1
767           1      93           70           31         0  30.4
```

| | DiabetesPedigreeFunction | Age | Outcome |
|-----|--------------------------|-----|---------|
| 763 | 0.171 | 63 | 0 |
| 764 | 0.340 | 27 | 0 |
| 765 | 0.245 | 30 | 0 |
| 766 | 0.349 | 47 | 1 |
| 767 | 0.315 | 23 | 0 |

```
[10]: data.isnull().any()
```

```
[10]: Pregnancies      False
      Glucose         False
      BloodPressure   False
      SkinThickness   False
      Insulin         False
      BMI            False
      DiabetesPedigreeFunction False
      Age            False
      Outcome        False
      dtype: bool
```

```
[11]: Positive = data[data['Outcome']==1]
      Positive.head(5)
```

```
[11]:      Pregnancies  Glucose  BloodPressure  SkinThickness  Insulin   BMI  \
0           6     148           72           35         0  33.6
2           8     183           64           0         0  23.3
4           0     137           40           35       168  43.1
6           3      78           50           32        88  31.0
8           2     197           70           45       543  30.5
```


| | DiabetesPedigreeFunction | Age | Outcome |
|---|--------------------------|-----|---------|
| 0 | 0.627 | 50 | 1 |
| 2 | 0.672 | 32 | 1 |
| 4 | 2.288 | 33 | 1 |
| 6 | 0.248 | 26 | 1 |
| 8 | 0.158 | 53 | 1 |

```
[12]: data['Glucose'].value_counts().head(7)
```

```
[12]: 99      17
      100     17
      111     14
      129     14
      125     14
      106     14
      112     13
      Name: Glucose, dtype: int64
```

```
[13]: plt.hist(data['Glucose'])
```

```
[13]: (array([ 5.,  0.,  4., 32., 156., 211., 163., 95., 56., 46.]),
      array([ 0., 19.9, 39.8, 59.7, 79.6, 99.5, 119.4, 139.3, 159.2,
              179.1, 199. ]),
      <a list of 10 Patch objects>)
```




```
[14]: data['BloodPressure'].value_counts().head(7)
```

```
[14]: 70    57
      74    52
      78    45
      68    45
      72    44
      64    43
      80    40
      Name: BloodPressure, dtype: int64
```

```
[15]: plt.hist(data['BloodPressure'])
```

```
[15]: (array([ 35.,   1.,   2.,  13., 107., 261., 243.,  87.,  14.,   5.]),
      array([ 0. , 12.2, 24.4, 36.6, 48.8, 61. , 73.2, 85.4, 97.6,
            109.8, 122. ]),
      <a list of 10 Patch objects>)
```




```
[16]: data['SkinThickness'].value_counts().head(7)
```

```
[16]: 0      227
      32      31
      30      27
      27      23
      23      22
      33      20
      28      20
      Name: SkinThickness, dtype: int64
```

```
[17]: plt.hist(data['SkinThickness'])
```

```
[17]: (array([231., 107., 165., 175., 78., 9., 2., 0., 0., 1.]),  
      array([ 0. , 9.9, 19.8, 29.7, 39.6, 49.5, 59.4, 69.3, 79.2, 89.1, 99. ]),  
      <a list of 10 Patch objects>)
```




```
[18]: data['Insulin'].value_counts().head(7)
```

```
[18]: 0      374  
      105     11  
      130      9  
      140      9  
      120      8  
      94       7  
      180       7  
      Name: Insulin, dtype: int64
```

```
[19]: plt.hist(data['Insulin'])
```

```
[19]: (array([487., 155., 70., 30., 8., 9., 5., 1., 2., 1.]),  
      array([ 0. , 84.6, 169.2, 253.8, 338.4, 423. , 507.6, 592.2, 676.8,  
              761.4, 846. ]),  
      <a list of 10 Patch objects>)
```




```
[20]: data['BMI'].value_counts().head(7)
```

```
[20]: 32.0    13
      31.6    12
      31.2    12
      0.0    11
      32.4    10
      33.3    10
      30.1     9
      Name: BMI, dtype: int64
```

```
[21]: plt.hist(data['BMI'])
```

```
[21]: (array([ 11.,  0., 15., 156., 268., 224., 78., 12.,  3.,  1.]),
      array([ 0. ,  6.71, 13.42, 20.13, 26.84, 33.55, 40.26, 46.97, 53.68,
              60.39, 67.1 ]),
      <a list of 10 Patch objects>)
```



```
[22]: data.describe().transpose()
```

```
[22]:
```


| | count | mean | std | min | 25% | \ |
|--------------------------|-------|------------|------------|--------|----------|---|
| Pregnancies | 768.0 | 3.845052 | 3.369578 | 0.000 | 1.00000 | |
| Glucose | 768.0 | 120.894531 | 31.972618 | 0.000 | 99.00000 | |
| BloodPressure | 768.0 | 69.105469 | 19.355807 | 0.000 | 62.00000 | |
| SkinThickness | 768.0 | 20.536458 | 15.952218 | 0.000 | 0.00000 | |
| Insulin | 768.0 | 79.799479 | 115.244002 | 0.000 | 0.00000 | |
| BMI | 768.0 | 31.992578 | 7.884160 | 0.000 | 27.30000 | |
| DiabetesPedigreeFunction | 768.0 | 0.471876 | 0.331329 | 0.078 | 0.24375 | |
| Age | 768.0 | 33.240885 | 11.760232 | 21.000 | 24.00000 | |
| Outcome | 768.0 | 0.348958 | 0.476951 | 0.000 | 0.00000 | |

| | 50% | 75% | max |
|--------------------------|----------|-----------|--------|
| Pregnancies | 3.0000 | 6.00000 | 17.00 |
| Glucose | 117.0000 | 140.25000 | 199.00 |
| BloodPressure | 72.0000 | 80.00000 | 122.00 |
| SkinThickness | 23.0000 | 32.00000 | 99.00 |
| Insulin | 30.5000 | 127.25000 | 846.00 |
| BMI | 32.0000 | 36.60000 | 67.10 |
| DiabetesPedigreeFunction | 0.3725 | 0.62625 | 2.42 |
| Age | 29.0000 | 41.00000 | 81.00 |
| Outcome | 0.0000 | 1.00000 | 1.00 |

1 Week:-2

```
[23]: plt.hist(Positive['BMI'],histtype='stepfilled',bins=20)
```

```
[23]: (array([ 2.,  0.,  0.,  0.,  0.,  0.,  3., 13., 38., 61., 61., 36., 27.,
          14.,  7.,  3.,  1.,  1.,  0.,  1.]),
      array([ 0.    ,  3.355,  6.71 , 10.065, 13.42 , 16.775, 20.13 , 23.485,
          26.84 , 30.195, 33.55 , 36.905, 40.26 , 43.615, 46.97 , 50.325,
          53.68 , 57.035, 60.39 , 63.745, 67.1  ]),
      <a list of 1 Patch objects>)
```




```
[24]: Positive['BMI'].value_counts().head(7)
```

```
[24]: 32.9      8
      31.6      7
      33.3      6
      31.2      5
      30.5      5
      32.0      5
      34.3      4
      Name: BMI, dtype: int64
```

```
[25]: plt.hist(Positive['Glucose'],histtype='stepfilled',bins=20)
```

```
[25]: (array([ 2.,  0.,  0.,  0.,  0.,  0.,  0.,  1.,  4.,  9., 28., 26., 36.,
            27., 29., 22., 24., 21., 25., 14.]),
      array([ 0. ,  9.95, 19.9 , 29.85, 39.8 , 49.75, 59.7 , 69.65,
            79.6 , 89.55, 99.5 , 109.45, 119.4 , 129.35, 139.3 , 149.25,
            159.2 , 169.15, 179.1 , 189.05, 199.  ]),
      <a list of 1 Patch objects>)
```




```
[26]: Positive['Glucose'].value_counts().head(7)
```

```
[26]: 125      7
      128      6
      129      6
      115      6
      158      6
      146      5
      124      5
      Name: Glucose, dtype: int64
```

```
[27]: plt.hist(Positive['BloodPressure'],histtype='stepfilled',bins=20)
```

```
[27]: (array([16.,  0.,  0.,  0.,  0.,  1.,  0.,  1.,  6.,  6., 19., 37., 56.,
            36., 41., 31.,  7.,  4.,  4.,  3.]),
      array([ 0. ,  5.7, 11.4, 17.1, 22.8, 28.5, 34.2, 39.9, 45.6,
            51.3, 57. , 62.7, 68.4, 74.1, 79.8, 85.5, 91.2, 96.9,
            102.6, 108.3, 114. ]),
      <a list of 1 Patch objects>)
```




```
[28]: Positive['BloodPressure'].value_counts().head(7)
```

```
[28]: 70    23
      76    18
      78    17
      74    17
      72    16
      0    16
      80    13
      Name: BloodPressure, dtype: int64
```

```
[29]: plt.hist(Positive['SkinThickness'],histtype='stepfilled',bins=20)
```

```
[29]: (array([88.,  1.,  4., 10., 18., 30., 41., 34., 23., 15.,  1.,  1.,  1.,
          0.,  0.,  0.,  0.,  0.,  0.,  1.]),
      array([ 0. ,  4.95,  9.9 , 14.85, 19.8 , 24.75, 29.7 , 34.65, 39.6 ,
          44.55, 49.5 , 54.45, 59.4 , 64.35, 69.3 , 74.25, 79.2 , 84.15,
          89.1 , 94.05, 99.  ]),
      <a list of 1 Patch objects>)
```




```
[30]: Positive['SkinThickness'].value_counts().head(7)
```

```
[30]: 0      88
      32     14
      30      9
      33      9
      39      8
      37      8
      36      8
      Name: SkinThickness, dtype: int64
```

```
[31]: plt.hist(Positive['Insulin'],histtype='stepfilled',bins=20)
```

```
[31]: (array([141.,  6., 23., 33., 24., 12.,  7.,  7.,  2.,  1.,  1.,
           5.,  3.,  1.,  1.,  0.,  0.,  0.,  0.,  1.]),
      array([ 0. , 42.3, 84.6, 126.9, 169.2, 211.5, 253.8, 296.1, 338.4,
           380.7, 423. , 465.3, 507.6, 549.9, 592.2, 634.5, 676.8, 719.1,
           761.4, 803.7, 846. ]),
      <a list of 1 Patch objects>)
```




```
[32]: Positive['Insulin'].value_counts().head(7)
```

```
[32]: 0      138
      130      6
      180      4
      175      3
      156      3
      185      2
      194      2
      Name: Insulin, dtype: int64
```


2 Scatter plot

```
[33]: BloodPressure = Positive['BloodPressure']
      Glucose = Positive['Glucose']
      SkinThickness = Positive['SkinThickness']
      Insulin = Positive['Insulin']
      BMI = Positive['BMI']
```


```
[34]: plt.scatter(BloodPressure, Glucose, color=['b'])
      plt.xlabel('BloodPressure')
      plt.ylabel('Glucose')
      plt.title('BloodPressure & Glucose')
      plt.show()
```




```
[37]: g =sns.scatterplot(x= "Glucose" ,y= "BloodPressure",  
                        hue="Outcome",  
                        data=data);
```



```
[38]: B =sns.scatterplot(x= "BMI" ,y= "Insulin",  
                        hue="Outcome",  
                        data=data);
```



```
[39]: S =sns.scatterplot(x= "SkinThickness" ,y= "Insulin",  
                        hue="Outcome",  
                        data=data);
```



Correlation Matrix

```
[40]: data.corr()
```

```
[40]:
```

| | Pregnancies | Glucose | BloodPressure | SkinThickness | \ |
|--------------------------|-------------|----------|---------------|---------------|---|
| Pregnancies | 1.000000 | 0.129459 | 0.141282 | -0.081672 | |
| Glucose | 0.129459 | 1.000000 | 0.152590 | 0.057328 | |
| BloodPressure | 0.141282 | 0.152590 | 1.000000 | 0.207371 | |
| SkinThickness | -0.081672 | 0.057328 | 0.207371 | 1.000000 | |
| Insulin | -0.073535 | 0.331357 | 0.088933 | 0.436783 | |
| BMI | 0.017683 | 0.221071 | 0.281805 | 0.392573 | |
| DiabetesPedigreeFunction | -0.033523 | 0.137337 | 0.041265 | 0.183928 | |
| Age | 0.544341 | 0.263514 | 0.239528 | -0.113970 | |
| Outcome | 0.221898 | 0.466581 | 0.065068 | 0.074752 | |


| | Insulin | BMI | DiabetesPedigreeFunction | \ |
|--------------------------|-----------|----------|--------------------------|---|
| Pregnancies | -0.073535 | 0.017683 | -0.033523 | |
| Glucose | 0.331357 | 0.221071 | 0.137337 | |
| BloodPressure | 0.088933 | 0.281805 | 0.041265 | |
| SkinThickness | 0.436783 | 0.392573 | 0.183928 | |
| Insulin | 1.000000 | 0.197859 | 0.185071 | |
| BMI | 0.197859 | 1.000000 | 0.140647 | |
| DiabetesPedigreeFunction | 0.185071 | 0.140647 | 1.000000 | |
| Age | -0.042163 | 0.036242 | 0.033561 | |
| Outcome | 0.130548 | 0.292695 | 0.173844 | |

| | Age | Outcome |
|--------------------------|-----------|----------|
| Pregnancies | 0.544341 | 0.221898 |
| Glucose | 0.263514 | 0.466581 |
| BloodPressure | 0.239528 | 0.065068 |
| SkinThickness | -0.113970 | 0.074752 |
| Insulin | -0.042163 | 0.130548 |
| BMI | 0.036242 | 0.292695 |
| DiabetesPedigreeFunction | 0.033561 | 0.173844 |
| Age | 1.000000 | 0.238356 |
| Outcome | 0.238356 | 1.000000 |

create correlation heat map


```
[41]: sns.heatmap(data.corr())
```

```
[41]: <matplotlib.axes._subplots.AxesSubplot at 0x7f48a7d528d0>
```



```
[42]: plt.subplots(figsize=(8,8))
sns.heatmap(data.corr(),annot=True,cmap='viridis')
```

[42]: <matplotlib.axes._subplots.AxesSubplot at 0x7f48a7cf0ed0>



```
[43]: plt.subplots(figsize=(8,8))
sns.heatmap(data.corr(),annot=True)
```

[43]: <matplotlib.axes._subplots.AxesSubplot at 0x7f48a5496850>



```
[44]: # Logistic Regreation and model building
```

```
[45]: data.head(5)
```

```
[45]:   Pregnancies  Glucose  BloodPressure  SkinThickness  Insulin  BMI  \
0           6     148           72           35         0  33.6
1           1      85           66           29         0  26.6
2           8     183           64            0         0  23.3
3           1      89           66           23        94  28.1
4           0     137           40           35       168  43.1

   DiabetesPedigreeFunction  Age  Outcome
```

| | | | |
|---|-------|----|---|
| 0 | 0.627 | 50 | 1 |
| 1 | 0.351 | 31 | 0 |
| 2 | 0.672 | 32 | 1 |
| 3 | 0.167 | 21 | 0 |
| 4 | 2.288 | 33 | 1 |

```
[78]: features = data.iloc[:,[0,1,2,3,4,5,6,7]].values
      label = data.iloc[:,8].values
```

```
[79]: #Train test split
      from sklearn.model_selection import train_test_split
      X_train,X_test,y_train,y_test = train_test_split(features,
                                                         label,
                                                         test_size=0.2,
                                                         random_state =10)
```

```
[82]: #Create model
      from sklearn.linear_model import LogisticRegression
      model = LogisticRegression()
      model.fit(X_train,y_train)
```

```
[82]: LogisticRegression()
```

```
[83]: y_predict=model.predict(X_test)
```

```
[53]: y_predict
```

```
[53]: array([1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 1, 0, 0, 1, 0, 0,
            0, 0, 0, 1, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0,
            0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
            1, 0, 0, 1, 1, 0, 0, 1, 0, 0, 1, 1, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0,
            1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 1, 0, 0, 0, 1, 1, 0, 1,
            0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0,
            1, 0, 0, 0, 0, 1, 1, 1, 0, 0, 0, 0, 1, 0, 0, 1, 1, 0, 0, 0, 1, 0])
```

```
[84]: print(model.score(X_train,y_train))
      print(model.score(X_test,y_test))
```

```
0.7719869706840391
0.7662337662337663
```

```
[85]: from sklearn.metrics import confusion_matrix
      cm = confusion_matrix(label,model.predict(features))
      cm
```

```
[85]: array([[446,  54],
            [122, 146]])
```

```
[86]: from sklearn.metrics import classification_report
print(classification_report(label,model.predict(features)))
```


| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.79 | 0.89 | 0.84 | 500 |
| 1 | 0.73 | 0.54 | 0.62 | 268 |
| accuracy | | | 0.77 | 768 |
| macro avg | 0.76 | 0.72 | 0.73 | 768 |
| weighted avg | 0.77 | 0.77 | 0.76 | 768 |

```
[87]: #Preparing ROC Curve (Receiver Operating Characteristics Curve)
from sklearn.metrics import roc_curve
from sklearn.metrics import roc_auc_score
```

```
[88]: # predict probabilities
probs = model.predict_proba(features)
# keep probabilities for the positive outcome only
probs = probs[:, 1]
# calculate AUC
auc = roc_auc_score(label, probs)
print('AUC: %.3f' % auc)
# calculate roc curve
fpr, tpr, thresholds = roc_curve(label, probs)
# plot no skill
plt.plot([0, 1], [0, 1], linestyle='--')
# plot the roc curve for the model
plt.plot(fpr, tpr, marker='.')
```

AUC: 0.837

```
[88]: [<matplotlib.lines.Line2D at 0x7f4842f5d910>]
```



```
[89]: #Applying Decission Tree Classifier
from sklearn.tree import DecisionTreeClassifier
model3 = DecisionTreeClassifier(max_depth=5)
model3.fit(X_train,y_train)
```

```
[89]: DecisionTreeClassifier(max_depth=5)
```

```
[90]: model3.score(X_train,y_train)
```

```
[90]: 0.8289902280130294
```

```
[91]: model3.score(X_test,y_test)
```

```
[91]: 0.7662337662337663
```

```
[92]: #Applying Random Forest
from sklearn.ensemble import RandomForestClassifier
model4 = RandomForestClassifier(n_estimators=11)
model4.fit(X_train,y_train)
```

```
[92]: RandomForestClassifier(n_estimators=11)
```

```
[93]: model4.score(X_train,y_train)
```

```
[93]: 0.99185667752443
```

```
[94]: model4.score(X_test,y_test)
```

```
[94]: 0.7467532467532467
```

```
[95]: #Support Vector Classifier  
  
from sklearn.svm import SVC
```

```
[96]: model5 = SVC(kernel='rbf',  
                 gamma='auto')  
model5.fit(X_train,y_train)
```

```
[96]: SVC(gamma='auto')
```

```
[99]: model5.score(X_test,y_test)
```

```
[99]: 0.6168831168831169
```

```
[100]: #Applying K-NN  
from sklearn.neighbors import KNeighborsClassifier  
model2 = KNeighborsClassifier(n_neighbors=7,  
                             metric='minkowski',  
                             p = 2)  
model2.fit(X_train,y_train)
```

```
[100]: KNeighborsClassifier(n_neighbors=7)
```

```
[101]: #Preparing ROC Curve (Receiver Operating Characteristics Curve)  
from sklearn.metrics import roc_curve  
from sklearn.metrics import roc_auc_score
```


```
[102]: # predict probabilities  
probs = model2.predict_proba(features)  
# keep probabilities for the positive outcome only  
probs = probs[:, 1]  
# calculate AUC  
auc = roc_auc_score(label, probs)  
print('AUC: %.3f' % auc)  
# calculate roc curve  
fpr, tpr, thresholds = roc_curve(label, probs)  
print("True Positive Rate - {}, False Positive Rate - {} Thresholds - {}".  
      ↪format(tpr,fpr,thresholds))  
# plot no skill  
plt.plot([0, 1], [0, 1], linestyle='--')  
# plot the roc curve for the model  
plt.plot(fpr, tpr, marker='.')  
plt.xlabel("False Positive Rate")
```

```
plt.ylabel("True Positive Rate")
```

AUC: 0.836

```
True Positive Rate - [0.          0.06716418 0.23880597 0.44776119 0.60074627
0.75373134
0.88059701 0.98507463 1.          ], False Positive Rate - [0.    0.    0.02
0.056 0.12  0.248 0.428 0.668 1.    ] Thresholds - [2.          1.
0.85714286 0.71428571 0.57142857 0.42857143
0.28571429 0.14285714 0.          ]
```

```
[102]: Text(0, 0.5, 'True Positive Rate')
```



```
[103]: #Precision Recall Curve for Logistic Regression
```

```
from sklearn.metrics import precision_recall_curve
from sklearn.metrics import f1_score
from sklearn.metrics import auc
from sklearn.metrics import average_precision_score
```

```
[104]: # predict probabilities
```

```
probs = model.predict_proba(features)
# keep probabilities for the positive outcome only
probs = probs[:, 1]
# predict class values
yhat = model.predict(features)
```




```

# calculate precision-recall curve
precision, recall, thresholds = precision_recall_curve(label, probs)
# calculate F1 score
f1 = f1_score(label, yhat)
# calculate precision-recall AUC
auc = auc(recall, precision)
# calculate average precision score
ap = average_precision_score(label, probs)
print('f1=%.3f auc=%.3f ap=%.3f' % (f1, auc, ap))
# plot no skill
plt.plot([0, 1], [0.5, 0.5], linestyle='--')
# plot the precision-recall curve for the model
plt.plot(recall, precision, marker='.')

```

f1=0.624 auc=0.726 ap=0.727

[104]: [<matplotlib.lines.Line2D at 0x7f4842e99090>]



```

[105]: #Precision Recall Curve for KNN


from sklearn.metrics import precision_recall_curve
from sklearn.metrics import f1_score
from sklearn.metrics import auc
from sklearn.metrics import average_precision_score

```

```
[106]: # predict probabilities
probs = model2.predict_proba(features)
# keep probabilities for the positive outcome only
probs = probs[:, 1]
# predict class values
yhat = model2.predict(features)
# calculate precision-recall curve
precision, recall, thresholds = precision_recall_curve(label, probs)
# calculate F1 score
f1 = f1_score(label, yhat)
# calculate precision-recall AUC
auc = auc(recall, precision)
# calculate average precision score
ap = average_precision_score(label, probs)
print('f1=%.3f auc=%.3f ap=%.3f' % (f1, auc, ap))
# plot no skill
plt.plot([0, 1], [0.5, 0.5], linestyle='--')
# plot the precision-recall curve for the model
plt.plot(recall, precision, marker='.'))
```

f1=0.658 auc=0.752 ap=0.709

[106]: [<matplotlib.lines.Line2D at 0x7f4842d40f50>]



[107]: #Precision Recall Curve for Decission Tree Classifier

```

from sklearn.metrics import precision_recall_curve
from sklearn.metrics import f1_score
from sklearn.metrics import auc
from sklearn.metrics import average_precision_score

```


```

[108]: # predict probabilities
probs = model3.predict_proba(features)
# keep probabilities for the positive outcome only
probs = probs[:, 1]
# predict class values
yhat = model3.predict(features)
# calculate precision-recall curve
precision, recall, thresholds = precision_recall_curve(label, probs)
# calculate F1 score
f1 = f1_score(label, yhat)
# calculate precision-recall AUC
auc = auc(recall, precision)
# calculate average precision score
ap = average_precision_score(label, probs)
print('f1=%.3f auc=%.3f ap=%.3f' % (f1, auc, ap))
# plot no skill
plt.plot([0, 1], [0.5, 0.5], linestyle='--')
# plot the precision-recall curve for the model
plt.plot(recall, precision, marker='.')

```

f1=0.709 auc=0.814 ap=0.769

[108]: [<matplotlib.lines.Line2D at 0x7f4842cb57d0>]



```
[109]: #Precision Recall Curve for Random Forest

from sklearn.metrics import precision_recall_curve
from sklearn.metrics import f1_score
from sklearn.metrics import auc
from sklearn.metrics import average_precision_score
```

```
[110]: # predict probabilities
probs = model4.predict_proba(features)
# keep probabilities for the positive outcome only
probs = probs[:, 1]
# predict class values
yhat = model4.predict(features)
# calculate precision-recall curve
precision, recall, thresholds = precision_recall_curve(label, probs)
# calculate F1 score
f1 = f1_score(label, yhat)
# calculate precision-recall AUC
auc = auc(recall, precision)
# calculate average precision score
ap = average_precision_score(label, probs)
print('f1=%.3f auc=%.3f ap=%.3f' % (f1, auc, ap))
# plot no skill
plt.plot([0, 1], [0.5, 0.5], linestyle='--')
# plot the precision-recall curve for the model
plt.plot(recall, precision, marker='.')
```

f1=0.915 auc=0.964 ap=0.957

```
[110]: [<matplotlib.lines.Line2D at 0x7f4842c2c050>]
```

